AEH - Advanced Application Engineer

Batch Eligible – 2026 Pass outs
Management Level – 11

ROLE OVERVIEW:

Work with diverse groups of innovators, to solve industries' and organizations' most challenging problems. Use leading-edge technologies and strategies to reinvent the way every job is done and help improve the way the world works and lives.
Create reusable software IP, crafting Small Language Models (SLMs) tailored to domain needs, and reimagining legacy workflows as autonomous, agentic systems.
As an early-career AI-native engineer (0–1 years of experience), you will be at the forefront of building domain-specific software platforms powered by GenAI and agentic codebases that drive IP-led services. Your mission is to co-create intelligent, self-evolving, and continuously learning systems that deliver transformational outcomes—at scale.

Key Responsibilities:
- Build Domain IP: Contribute to the creation of SLMs and reusable agentic modules tuned to specific industries (e.g., banking, healthcare, supply chain) that serve as the foundation of next-gen service offerings.
- Engineer with Purpose: Develop and evolve domain-aware AI agents that perform business tasks traditionally done by human workflows, from document understanding to decision-making.
- Architect the Software of the Future: Move beyond application delivery to design domain-centered code frameworks that are self-learning, interpretable, and modular—ready to be reused, extended, and monetized.
- AI-First Development: Use agentic development environments and AI-native tooling (e.g., Copilot, Claude, Gemini) to ideate, generate, test, and deploy software IP faster and smarter.
- Collaborate Across Domains: Work with domain experts, cognitive architects, and data scientists to embed enterprise knowledge into software, shaping how services will be delivered in an autonomous future.
- Reimagine Workflows: Help dismantle legacy system boundaries by re-orchestrating traditional processes into GenAI-powered experiences—infused with automation, prediction, and learning.
- Lead with Data and Models: Actively engage in curating domain datasets, fine-tuning SLMs, and experimenting with prompt engineering and retrieval-augmented generation to maximize outcome accuracy and relevance.
- Operationalize Innovation: Integrate your work into production using scalable cloud-native infrastructure, containerization, and MLOps/DevOps practices.
- Document technical specifications, maintain cloud resource diagrams, and keep track of development and deployment activities.
- Continuously explore and experiment with the latest in AI and cloud technologies to drive innovation and productivity improvements.

What we need?
We would love you to have:
Required Skills & Qualifications:

- Bachelor's or Master's in Computer Science, AI, or a related field.
- Proficiency in modern programming (e.g., Python, Java, Go, etc.) and exposure to LLM/SLM tooling.
- Awareness of transformer models, GenAI architectures, and frameworks like LangChain, RAG, or OpenLLM.
- Ability to think modular, reusable, and product-centric—building for IP, not just project delivery.
- Exposure to agentic AI platforms and tools (e.g., CrewAI, AutoGen, AutoGPT, OpenAgents).
- Familiarity with cloud-native technologies (AWS/GCP/Azure), serverless computing, and containerization.
- Strong analytical mindset, curiosity, and a passion for reshaping software with intelligence.

Preferred Skills:

- Experience building or fine-tuning SLMs or narrow-domain models.
- Exposure to building vector databases, semantic retrieval pipelines, and prompt engineering techniques.
- Understanding of domain modeling, ontology creation, and enterprise knowledge graphs.
- Knowledge of responsible AI principles and explainable AI practices.
- Early-stage contributions to open-source GenAI ecosystems or agent frameworks.

Accenture has not authorized any agency, company or individual to either collect money or arrive on any monetary arrangement in exchange for a job at Accenture.
Accenture's criterion for hiring candidates is merit. Any agency, company or individual offering employment with Accenture in exchange for money is misrepresenting their relationship with Accenture, which has not authorized any such action. If you are approached by any entity or individuals who demand money or any other form of compensation in return for a job offer at Accenture – even if they present themselves as representatives or employees of Accenture – please send the details to
Accenture Business Ethics Line https://businessethicsline.com/accenture